

zpfvw2obv

September 16, 2024

1 ANN Model for Early Warning System for Customer Support

```
[171]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
[172]: #reading the dataset
data = pd.read_csv('/content/drive/MyDrive/TechConsulting/Early Warning System_
↳for Customer Support/Churn_Modelling.csv')
data= data.drop(['RowNumber', 'CustomerId', 'Surname'],axis=1)
data
```

```
[172]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	\
0	619	France	Female	42	2	0.00	1	
1	608	Spain	Female	41	1	83807.86	1	
2	502	France	Female	42	8	159660.80	3	
3	699	France	Female	39	1	0.00	2	
4	850	Spain	Female	43	2	125510.82	1	
...	
9995	771	France	Male	39	5	0.00	2	
9996	516	France	Male	35	10	57369.61	1	
9997	709	France	Female	36	7	0.00	1	
9998	772	Germany	Male	42	3	75075.31	2	
9999	792	France	Female	28	4	130142.79	1	

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	1	101348.88	1
1	0	1	112542.58	0
2	1	0	113931.57	1
3	0	0	93826.63	0
4	1	1	79084.10	0
...
9995	1	0	96270.64	0
9996	1	1	101699.77	0
9997	0	1	42085.58	1
9998	1	0	92888.52	1
9999	1	0	38190.78	0

[10000 rows x 11 columns]

```
[173]: # data info
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CreditScore            10000 non-null  int64
1   Geography              10000 non-null  object
2   Gender                 10000 non-null  object
3   Age                    10000 non-null  int64
4   Tenure                 10000 non-null  int64
5   Balance                10000 non-null  float64
6   NumOfProducts          10000 non-null  int64
7   HasCrCard              10000 non-null  int64
8   IsActiveMember         10000 non-null  int64
9   EstimatedSalary        10000 non-null  float64
10  Exited                 10000 non-null  int64
dtypes: float64(2), int64(7), object(2)
memory usage: 859.5+ KB
```

```
[174]: # data description
data.describe()
```

```
[174]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	\
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	
mean	650.528800	38.921800	5.012800	76485.889288	1.530200	
std	96.653299	10.487806	2.892174	62397.405202	0.581654	
min	350.000000	18.000000	0.000000	0.000000	1.000000	
25%	584.000000	32.000000	3.000000	0.000000	1.000000	
50%	652.000000	37.000000	5.000000	97198.540000	1.000000	
75%	718.000000	44.000000	7.000000	127644.240000	2.000000	
max	850.000000	92.000000	10.000000	250898.090000	4.000000	

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	0.70550	0.515100	100090.239881	0.203700
std	0.45584	0.499797	57510.492818	0.402769
min	0.00000	0.000000	11.580000	0.000000
25%	0.00000	0.000000	51002.110000	0.000000
50%	1.00000	1.000000	100193.915000	0.000000
75%	1.00000	1.000000	149388.247500	0.000000
max	1.00000	1.000000	199992.480000	1.000000


```
[175]: # Checking missing values
data.isnull().sum()
```

```
[175]: CreditScore      0
Geography      0
Gender         0
Age            0
Tenure         0
Balance        0
NumOfProducts  0
HasCrCard      0
IsActiveMember 0
EstimatedSalary 0
Exited         0
dtype: int64
```


1.1 EDA

```
[176]: ## Corelation matrix
import seaborn as sns
plt.figure(figsize=(10, 5))
c = data.corr(numeric_only=True)
sns.heatmap(c, cmap='BrBG', annot=True)
```

```
[176]: <Axes: >
```



```
[177]: # target variable distribution
exited_counts = data['Exited'].value_counts()
plt.figure(figsize=(6,4))
plt.bar(exited_counts.index, exited_counts.values, color='skyblue',
        edgecolor='black')
plt.title('Distribution of Churn (Exited)')
plt.xlabel('Exited (0 = No, 1 = Yes)')
plt.ylabel('Count')
plt.xticks([0, 1], ['No (0)', 'Yes (1)'])
plt.show()
```



```
[178]: # varibale distribution
# Function to plot variable distributions
def plot_variable_distributions(df):
    features = data.columns[:-1]
    plt.figure(figsize=(15, 15))

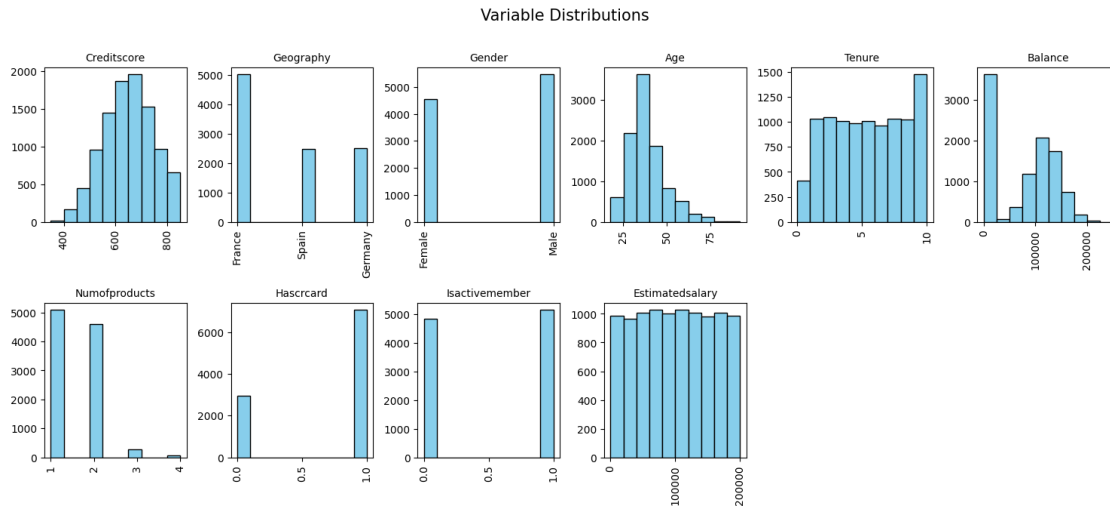
    for i, feature in enumerate(features, 1):
        plt.subplot(5, 6, i) # Adjust the grid size as per the number of
        features
        df[feature].hist(grid=False, color='skyblue', edgecolor='black')
        plt.title(feature.capitalize(), fontsize=10)
        plt.xlabel("")
```

```

plt.ylabel("")
plt.xticks(rotation=90)

plt.tight_layout(pad=1)
plt.suptitle("Variable Distributions", fontsize=15, y=1.03)
plt.show()
# Call the function to plot histograms of the DataFrame excluding the first
↪column
plot_variable_distributions(data)


```



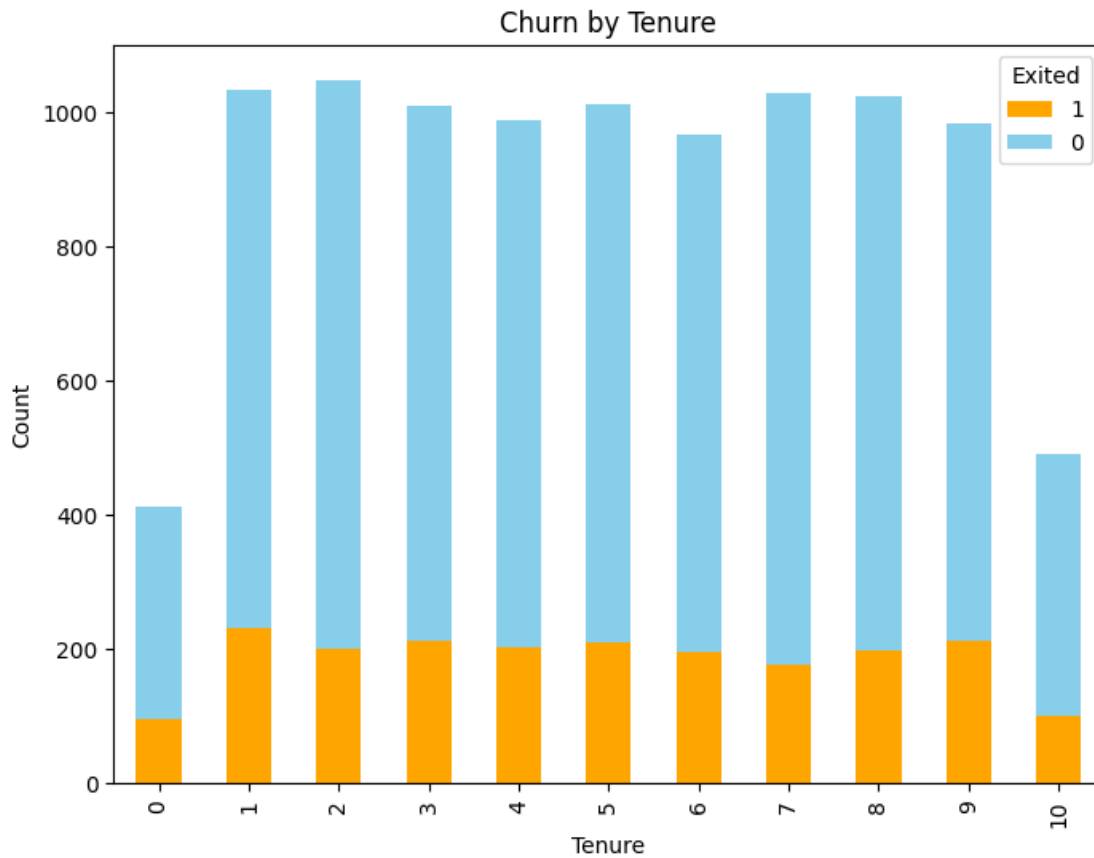
```

[179]: gender_churn = data.groupby(['Gender', 'Exited']).size().unstack()
gender_churn.plot(kind='bar', color=['skyblue', 'orange'], figsize=(6,4))
plt.title('Churn by Gender')
plt.xlabel('Gender')
plt.ylabel('Count')
plt.show()

```



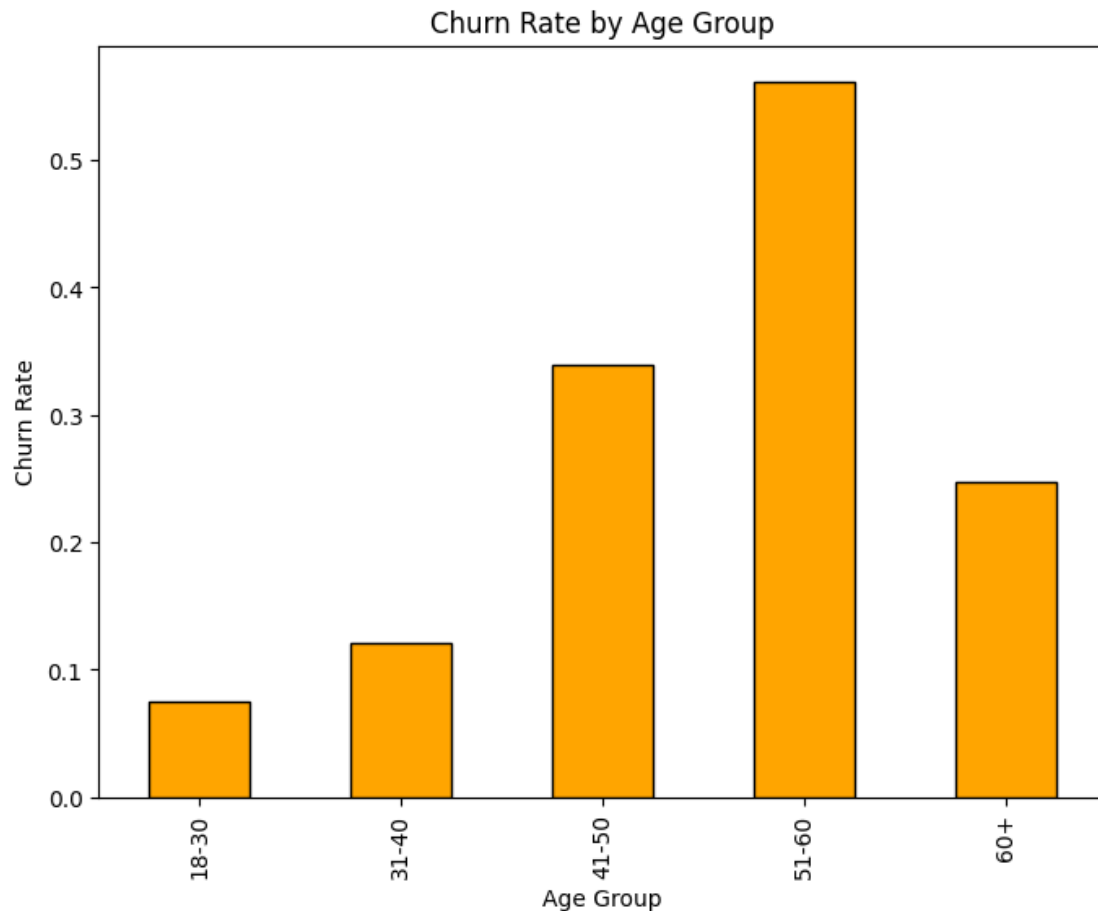
```
[180]: # Churn by Tenure
tenure_churn = data.groupby(['Tenure', 'Exited']).size().unstack()
tenure_churn = tenure_churn[[1, 0]]
tenure_churn.plot(kind='bar', stacked=True, color=['orange', 'skyblue'],
    figsize=(8,6))
plt.title('Churn by Tenure')
plt.xlabel('Tenure')
plt.ylabel('Count')
plt.show()
```



```
[181]: ## churn rate by Age
age_churn = data.groupby(pd.cut(data['Age'], bins=[18, 30, 40, 50, 60, 100],
    labels=['18-30', '31-40', '41-50', '51-60', '60+']))['Exited'].mean()
plt.figure(figsize=(8,6))
age_churn.plot(kind='bar', color='orange', edgecolor='black')
plt.title('Churn Rate by Age Group')
plt.xlabel('Age Group')
plt.ylabel('Churn Rate')
plt.show()
```

<ipython-input-181-ff4cd00f48ad>:2: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
age_churn = data.groupby(pd.cut(data['Age'], bins=[18, 30, 40, 50, 60, 100],
labels=['18-30', '31-40', '41-50', '51-60', '60+']))['Exited'].mean()
```




```
[182]: ## Churn by account balance
plt.figure(figsize=(8,6))
sns.boxplot(x='Exited', y='Balance', data=data, palette="Set2")
plt.title('Account Balance vs. Churn')
plt.xlabel('Churn (0 = No, 1 = Yes)')
plt.ylabel('Account Balance')
plt.show()
```


<ipython-input-182-8255b7de9b9a>:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.


```
sns.boxplot(x='Exited', y='Balance', data=data, palette="Set2")
```

```
[183]: ## Churn by geography
geography_churn = data.groupby('Geography')['Exited'].mean()
plt.figure(figsize=(8,6))
geography_churn.plot(kind='bar', color='lightblue', edgecolor='black')
plt.title('Churn Rate by Geography')
plt.xlabel('Geography')
plt.ylabel('Churn Rate')
plt.xticks([0, 1, 2], ['France', 'Spain', 'Germany'])
plt.show()
```



```
[184]: ## churn by number of products
plt.figure(figsize=(8,6))
sns.countplot(x='NumOfProducts', hue='Exited', data=data, palette='Set2')
plt.title('Churn by Number of Products')
plt.xlabel('Number of Products')
plt.ylabel('Count')
plt.show()
```



1.2 Preparation

[185]: data

```
[185]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	\
0	619	France	Female	42	2	0.00	1	
1	608	Spain	Female	41	1	83807.86	1	
2	502	France	Female	42	8	159660.80	3	
3	699	France	Female	39	1	0.00	2	
4	850	Spain	Female	43	2	125510.82	1	
...	
9995	771	France	Male	39	5	0.00	2	
9996	516	France	Male	35	10	57369.61	1	
9997	709	France	Female	36	7	0.00	1	
9998	772	Germany	Male	42	3	75075.31	2	
9999	792	France	Female	28	4	130142.79	1	

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	1	101348.88	1

1	0	1	112542.58	0
2	1	0	113931.57	1
3	0	0	93826.63	0
4	1	1	79084.10	0
...
9995	1	0	96270.64	0
9996	1	1	101699.77	0
9997	0	1	42085.58	1
9998	1	0	92888.52	1
9999	1	0	38190.78	0

[10000 rows x 11 columns]

```
[186]: ## encoding
data = pd.get_dummies(data, columns=['Geography', 'Gender'], drop_first=True,
dtype=int)
data
```

```
[186]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
0	619	42	2	0.00	1	1	
1	608	41	1	83807.86	1	0	
2	502	42	8	159660.80	3	1	
3	699	39	1	0.00	2	0	
4	850	43	2	125510.82	1	1	
...	
9995	771	39	5	0.00	2	1	
9996	516	35	10	57369.61	1	1	
9997	709	36	7	0.00	1	0	
9998	772	42	3	75075.31	2	1	
9999	792	28	4	130142.79	1	1	

	IsActiveMember	EstimatedSalary	Exited	Geography_Germany	\
0	1	101348.88	1	0	
1	1	112542.58	0	0	
2	0	113931.57	1	0	
3	0	93826.63	0	0	
4	1	79084.10	0	0	
...	
9995	0	96270.64	0	0	
9996	1	101699.77	0	0	
9997	1	42085.58	1	0	
9998	0	92888.52	1	1	
9999	0	38190.78	0	0	

	Geography_Spain	Gender_Male
0	0	0
1	1	0

2	0	0
3	0	0
4	1	0
...
9995	0	1
9996	0	1
9997	0	0
9998	0	1
9999	0	0

[10000 rows x 12 columns]

```
[187]: # ## Transformation
# data['CreditScore_log'] = np.log1p(data['CreditScore'])
# data['Age_log'] = np.log1p(data['Age'])
# data['Balance_log'] = np.log1p(data['Balance'])
# data['EstimatedSalary_log'] = np.log1p(data['EstimatedSalary'])
# data
```

```
[188]: data.shape
```

```
[188]: (10000, 12)
```

```
[189]: ## outliers
def treat_outliers(df, features):
    df_filtered = df.copy()
    for column in features:
        Q1 = df_filtered[column].quantile(0.25)
        Q3 = df_filtered[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        df_filtered = df_filtered[(df_filtered[column] >= lower_bound) &
        ↪ (df_filtered[column] <= upper_bound)]
    return df_filtered

n_features= ['CreditScore', 'Age', 'Balance', 'EstimatedSalary']
data_cleaned = treat_outliers(data, n_features)
data_cleaned.shape

#data_cleaned= data
```

```
[189]: (9626, 12)
```

```
[190]: data
```

```
[190]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
0	619	42	2	0.00	1	1	
1	608	41	1	83807.86	1	0	
2	502	42	8	159660.80	3	1	
3	699	39	1	0.00	2	0	
4	850	43	2	125510.82	1	1	
...	
9995	771	39	5	0.00	2	1	
9996	516	35	10	57369.61	1	1	
9997	709	36	7	0.00	1	0	
9998	772	42	3	75075.31	2	1	
9999	792	28	4	130142.79	1	1	

	IsActiveMember	EstimatedSalary	Exited	Geography_Germany	\
0	1	101348.88	1	0	
1	1	112542.58	0	0	
2	0	113931.57	1	0	
3	0	93826.63	0	0	
4	1	79084.10	0	0	
...	
9995	0	96270.64	0	0	
9996	1	101699.77	0	0	
9997	1	42085.58	1	0	
9998	0	92888.52	1	1	
9999	0	38190.78	0	0	

	Geography_Spain	Gender_Male
0	0	0
1	1	0
2	0	0
3	0	0
4	1	0
...
9995	0	1
9996	0	1
9997	0	0
9998	0	1
9999	0	0

[10000 rows x 12 columns]

```
[191]: ## Scaling
def min_max_scale(df, exclude_columns=None):
    df_scaled = df.copy()
    scaling_params = {}
    for column in df_scaled.columns:
        if column not in exclude_columns:
```

```

        min_value = df_scaled[column].min()
        max_value = df_scaled[column].max()
        df_scaled[column] = (df_scaled[column] - min_value) / (max_value -
↳min_value)

        scaling_params[column] = {'min': min_value, 'max': max_value}
    return df_scaled, scaling_params
exclude_columns = ['HasCrCard', 'IsActiveMember', 'Exited',
↳'Geography_Germany', 'Geography_Spain', 'Gender_Male']
data_scaled, scaling_params = min_max_scale(data_cleaned, exclude_columns)
data_scaled

```

```

[191]:      CreditScore      Age  Tenure  Balance  NumOfProducts  HasCrCard  \
0      0.505353  0.545455    0.2  0.000000    0.000000    1
1      0.481799  0.522727    0.1  0.334031    0.000000    0
2      0.254818  0.545455    0.8  0.636357    0.666667    1
3      0.676660  0.477273    0.1  0.000000    0.333333    0
4      1.000000  0.568182    0.2  0.500246    0.000000    1
...      ...      ...      ...      ...      ...      ...
9995     0.830835  0.477273    0.5  0.000000    0.333333    1
9996     0.284797  0.386364    1.0  0.228657    0.000000    1
9997     0.698073  0.409091    0.7  0.000000    0.000000    0
9998     0.832976  0.545455    0.3  0.299226    0.333333    1
9999     0.875803  0.227273    0.4  0.518708    0.000000    1

```

```

      IsActiveMember  EstimatedSalary  Exited  Geography_Germany  \
0      1      0.506735    1      0
1      1      0.562709    0      0
2      0      0.569654    1      0
3      0      0.469120    0      0
4      1      0.395400    0      0
...      ...      ...      ...      ...
9995     0      0.481341    0      0
9996     1      0.508490    0      0
9997     1      0.210390    1      0
9998     0      0.464429    1      1
9999     0      0.190914    0      0

```

```

      Geography_Spain  Gender_Male
0      0      0
1      1      0
2      0      0
3      0      0
4      1      0
...      ...      ...
9995     0      1
9996     0      1
9997     0      0

```

```

9998          0          1
9999          0          0

```

[9626 rows x 12 columns]

1.3 Modelling

```

[192]: ## data for modelling
model_data = data_scaled.copy()
model_data

```

```

[192]:      CreditScore      Age  Tenure  Balance  NumOfProducts  HasCrCard  \
0      0.505353  0.545455    0.2  0.000000    0.000000    1
1      0.481799  0.522727    0.1  0.334031    0.000000    0
2      0.254818  0.545455    0.8  0.636357    0.666667    1
3      0.676660  0.477273    0.1  0.000000    0.333333    0
4      1.000000  0.568182    0.2  0.500246    0.000000    1
...      ...      ...      ...      ...      ...      ...
9995     0.830835  0.477273    0.5  0.000000    0.333333    1
9996     0.284797  0.386364    1.0  0.228657    0.000000    1
9997     0.698073  0.409091    0.7  0.000000    0.000000    0
9998     0.832976  0.545455    0.3  0.299226    0.333333    1
9999     0.875803  0.227273    0.4  0.518708    0.000000    1

```

```

      IsActiveMember  EstimatedSalary  Exited  Geography_Germany  \
0                  1      0.506735      1      0
1                  1      0.562709      0      0
2                  0      0.569654      1      0
3                  0      0.469120      0      0
4                  1      0.395400      0      0
...      ...      ...      ...      ...
9995              0      0.481341      0      0
9996              1      0.508490      0      0
9997              1      0.210390      1      0
9998              0      0.464429      1      1
9999              0      0.190914      0      0

```

```

      Geography_Spain  Gender_Male
0                  0          0
1                  1          0
2                  0          0
3                  0          0
4                  1          0
...      ...      ...
9995              0          1
9996              0          1
9997              0          0

```



```

9998          0          1
9999          0          0

```

[9626 rows x 12 columns]

```

[193]: ## checking imbalance
model_data['Exited'].value_counts()

```

```

[193]: Exited
0      7677
1      1949
Name: count, dtype: int64

```

```

[194]: ## over sampling
majority_class = model_data[model_data['Exited'] == 0]
minority_class = model_data[model_data['Exited'] == 1]
minority_class_oversampled = minority_class.sample(len(majority_class),
    ↪replace=True, random_state=42)
oversampled_data = pd.concat([majority_class, minority_class_oversampled])
oversampled_data = oversampled_data.sample(frac=1, random_state=42).
    ↪reset_index(drop=True)
oversampled_data['Exited'].value_counts()

#oversampled_data = model_data

```

```

[194]: Exited
1      7677
0      7677
Name: count, dtype: int64

```

```

[195]: ## train test split
train_size = int(0.8 * len(oversampled_data))
train_df = oversampled_data[:train_size]
test_df = oversampled_data[train_size:]

X_train = train_df.drop('Exited', axis=1)
y_train = train_df['Exited']
X_test = test_df.drop('Exited', axis=1)
y_test = test_df['Exited']
len(X_train), len(y_train), len(X_test), len(y_test)

```

```

[195]: (12283, 12283, 3071, 3071)

```

```

[196]: ## to_numpy
X_train= X_train.to_numpy()
y_train= y_train.to_numpy()
X_test= X_test.to_numpy()

```

```
y_test= y_test.to_numpy()
set(y_train)
```

```
[196]: {0, 1}
```

2 Binary Logistic Regression Class

2.1 Useful Functions

```
[197]: #losses, activations, metrics
def sigmoid(h):
    return 1 / (1 + np.exp(-h))

def cross_entropy(y, p_hat):
    return -(1/len(y))*np.sum(y*np.log(p_hat)+ (1-y)*np.log(1-p_hat))

def accuracy(y, y_hat):
    return np.mean(y == y_hat)
```

```
[198]: class LogisticRegression():

    def __init__(self, thresh=0.5):
        self.thresh = thresh
        self.W = None
        self.b = None

    def fit(self, X, y, eta=1e-3, epochs = 1e3, show_curve=False):
        epochs = int(epochs)
        N, D = X.shape

        #Initialize Weight and biases
        self.W= np.random.randn(D)
        self.b= np.random.randn(1)
        #Create Zero values container for J
        J = np.zeros(epochs)

        #SGD
        for epoch in range(epochs):
            p_hat= self.__forward__(X)
            J[epoch]= cross_entropy(y, p_hat)

            #Weight and bias Update Rules
            self.W -= eta*(1/N)*X.T@(p_hat-y)
            self.b -= eta*(1/N)*np.sum(p_hat-y)

        if show_curve:
            plt.figure()
```

```

plt.plot(J)
plt.xlabel('epochs')
plt.ylabel('$\mathcal{J}$')
plt.title('Training Curve')
plt.show()

def __forward__(self, X):
    return sigmoid(X@self.W+self.b)

def predict(self, X):
    return (self.__forward__(X)>= self.thresh).astype(np.int32)

```

```
[199]: log_reg= LogisticRegression()
```

```
[136]: epochs_list = [3500, 4000, 4500]
eta_list = [0.08, 0.1, 0.12]

best_accuracy = 0
best_params = {'epochs': None, 'eta': None}

for epochs in epochs_list:
    for eta in eta_list:
        log_reg.fit(X_train, y_train, epochs=epochs, eta=eta, show_curve=False)

        # Predict and calculate accuracy
        y_hat_BLR = log_reg.predict(X_test)
        accuracy_val = accuracy(y_test, y_hat_BLR)

        # Track the best accuracy and parameters
        if accuracy_val > best_accuracy:
            best_accuracy = accuracy_val
            best_params['epochs'] = epochs
            best_params['eta'] = eta

print(f"Best Accuracy: {best_accuracy:.4f}")
print(f"Best Parameters: Epochs={best_params['epochs']},
↪Eta={best_params['eta']}")


```

Best Accuracy: 0.7258

Best Parameters: Epochs=4500, Eta=0.12

```
[200]: log_reg.fit(X_train, y_train, epochs= 4500, eta= 0.12, show_curve= True)
y_hat_BLR = log_reg.predict(X_test)
print(f"Training Accuracy: {accuracy(y_test, y_hat_BLR): 0.4f}")
print(log_reg.W)
print(log_reg.b)

```




Training Accuracy: 0.7245

```
[-0.31768667  4.52905624 -0.23857845  0.70784429 -0.33213243 -0.13377152
 -0.77718349 -0.04615254  0.85188081  0.02771341 -0.53021483]
[-1.68362384]
```

```
[201]: #Binary Logistic Regression Class
plt.figure(figsize=(10,7))
y_actual_BLR = pd.Series(y_test, name='Actual')
y_pred_BLR = pd.Series(y_hat_BLR, name='Predicted')
cm = pd.crosstab(y_actual_BLR, y_pred_BLR)
ax = sns.heatmap(cm, annot=True, fmt="d")
plt.title("Binary Logistic Regression Class Confusion Matrix")
plt.ylabel('True label')
plt.xlabel('Predicted label')
```

```
[201]: Text(0.5, 47.722222222222, 'Predicted label')
```



```
[138]: # #Binary Logistic Regression Class
# plt.figure(figsize=(10,7))
# y_actual_BLR = pd.Series(y_test, name='Actual')
# y_pred_BLR = pd.Series(y_hat_BLR, name='Predicted')
# cm = pd.crosstab(y_actual_BLR, y_pred_BLR)
# ax = sns.heatmap(cm, annot=True, fmt="d")
# plt.title("Binary Logistic Regression Class Confusion Matrix")
# plt.ylabel('True label')
# plt.xlabel('Predicted label')
```

3 Two-Layer Feed Forward Perceptron

```
[202]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

3.1 Activation Functions

```
[203]: def linear(H):  
        return H  
  
        def ReLU(H):  
            return H*(H>0)  
  
        def sigmoid(H):  
            return 1/(1+np.exp(-H))  
  
        def softmax(H):  
            eH = np.exp(H)  
            return eH/eH.sum(axis=1, keepdims= True)  
  
        #np.tanh
```

3.2 Useful Functions

```
[204]: def one_hot_encode(y):  
        N = len(y)  
        K =len(set(y))  
        Y = np.zeros((N,K))  
        for i in range(N):  
            Y[i, y[i]]=1  
        return Y  
  
        def cross_entropy(Y, P_hat):  
            return -np.sum(Y*np.log(P_hat))  
  
        def binary_cross_entropy(y, p_hat):  
            return -(1/len(y))*np.sum(y*np.log(p_hat)+ (1-y)*np.log(1-p_hat))  
  
        def accuracy(y, y_hat):  
            return np.mean(y == y_hat)
```

3.3 Shallow ANN Class

```
[205]: class Shallow_ANN():  
  
        def fit(self, X, y, neurons=6, eta=1e-3, epochs=1e3, show_curve= True):  
            epochs= int(epochs)  
            N, D = X.shape  
            Y = one_hot_encode(y)  
            #Y = y.reshape(-1, 1)  
            #Y=y  
            K= Y.shape[1]
```

```

#Weights Initialization
self.W = {l: np.random.randn(M[0], M[1]) for l, M in enumerate(zip([D, u
↪neurons], [neurons, K]), 1)}
self.B = {l: np.random.randn(M) for l, M in enumerate([neurons, K], 1)}
#Define Activations
self.a = {1: np.tanh, 2: softmax}

J = np.zeros(epochs)

#SGD Steps
for epoch in range(epochs):
    self.__forward__(X)
    J[epoch] = cross_entropy(Y, self.Z[2])
    #Weight update rules for output layer(Layer 2)
    self.W[2] -= eta*(1/N)* self.Z[1].T@(self.Z[2] - Y)
    self.B[2] -= eta*(1/N)*(self.Z[2]-Y).sum(axis=0)
    #Weight Update Rule for Layer 1
    #self.W[1] -= eta*(1/N)*X.T@((self.Z[2]-Y)@self.W[2].T*(1-self.Z[1]**2))
    #self.B[1] -= eta*(1/N)*((self.Z[2]-Y)@self.W[2].T*(1-self.Z[1]**2)).
    ↪sum(axis=0)
    self.W[1] -= eta*(1/N)*X.T@((self.Z[2]-Y)@self.W[2].T*(1-self.Z[1]**2))
    self.B[1] -= eta*(1/N)*((self.Z[2]-Y)@self.W[2].T*(1-self.Z[1]**2)).
    ↪sum(axis=0)

    if show_curve:
        plt.figure()
        plt.plot(J)
        plt.xlabel('epochs')
        plt.ylabel('$\mathcal{J}$')
        plt.title('Training Curve')
        plt.show()

def __forward__(self, X):
    self.Z = {0: X}
    for l in sorted(self.W.keys()):
        self.Z[l] = self.a[l](self.Z[l-1]@self.W[l] + self.B[l])

def predict(self, X):
    self.__forward__(X)
    return self.Z[2].argmax(axis=1)
    #return (self.Z[2] > 0.5).astype(int)

```

```

[209]: # Define ranges for hyperparameters to test
neurons_list = [2, 4, 6, 8] # Number of neurons to try
eta_list = [0.01, 0.05, 0.1] # Learning rates to test
epochs_list = [2000, 4000, 8000] # Epochs to test

```

```

best_accuracy = 0
best_params = {'neurons': None, 'eta': None, 'epochs': None}

# Loop through all combinations of hyperparameters
for neurons in neurons_list:
    for eta in eta_list:
        for epochs in epochs_list:
            print(f"Trying: Neurons={neurons}, Eta={eta}, Epochs={epochs}")

            # Initialize the model with the current configuration
            my_ann = Shallow_ANN()

            # Fit the model with current hyperparameters
            my_ann.fit(X_train, y_train, neurons=neurons, eta=eta,
↪epochs=epochs, show_curve=False)

            # Predict on the test set
            y_hat_2LFF = my_ann.predict(X_test)

            # Calculate accuracy
            accuracy_val = accuracy(y_test, y_hat_2LFF)
            print(f"Accuracy: {accuracy_val:.4f} | Neurons: {neurons}, Eta:
↪{eta}, Epochs: {epochs}")

            # Track the best performing hyperparameters
            if accuracy_val > best_accuracy:
                best_accuracy = accuracy_val
                best_params['neurons'] = neurons
                best_params['eta'] = eta
                best_params['epochs'] = epochs

# Print the best result
print(f"Best Accuracy: {best_accuracy:.4f}")
print(f"Best Parameters: Neurons={best_params['neurons']},
↪Eta={best_params['eta']}, Epochs={best_params['epochs']}")

```

```

[212]: my_ann = Shallow_ANN()
#neurons= 2, eta=0.001, epochs = 30000
#eta=0.1, epochs=4000
# Neurons=4, Eta=0.1, Epochs=8000
my_ann.fit(X_train, y_train, neurons= 4, eta=0.1, epochs = 8000, show_curve=
↪True )
#my_ann.fit(X_train, y_train, neurons= 6, eta=1e-3, epochs=1e4, show_curve=
↪True )
y_hat_2LFF =my_ann.predict(X_test)


```



```

print("Accuracy: ", accuracy(y_test, y_hat_2LFF))
print(my_ann.W)
print(my_ann.B)

```



```


Accuracy: 0.7323347443829371
{1: array([[ 0.62799951,  0.53056275, -1.13659122, -0.01950681],
          [-0.51098416, -1.42374138, -0.66142178,  1.98602998],
          [-0.11328237, -0.5119192 , -0.97500364, -0.09710407],
          [-0.97496391, -0.36327879,  0.02422441,  0.251694  ],
          [-2.17776282,  0.31109604,  0.27490795, -0.21974436],
          [ 0.08837398, -0.57468743,  0.59100345, -0.03834669],
          [-0.48482615, -1.23939884,  0.79512002, -0.34500277],
          [ 0.63977724, -1.26610834,  0.13822085,  0.02210972],
          [-0.25397988,  0.1944344 , -0.70004899,  0.33049197],
          [ 0.65995321, -0.11790284,  0.39217273,  0.0213824  ],
          [ 3.92258091, -0.56943913, -1.26083242,  0.13298094]]), 2: array([[
0.05297944, -0.44038011],
          [ 0.17970899, -0.08381809],
          [ 0.62186464,  0.41513336],
          [-2.55111934,  0.43739169]]))}

```

```
{1: array([-2.09331761, -0.57821027, -1.58123866, -1.31860084]), 2:
array([0.62954312, 0.95837611])}
```

```
[213]: #.reshape(-1)
## Two-Layer Feed Forward Perceptron
plt.figure(figsize=(10,7))
y_actual_2LFF = pd.Series(y_test, name='Actual')
y_pred_2LFF = pd.Series(y_hat_2LFF, name='Predicted')
cm = pd.crosstab(y_actual_2LFF, y_pred_2LFF)
ax = sns.heatmap(cm, annot=True, fmt="d")
plt.title("Two-Layer Feed Forward Perceptron Confusion Matrix")
plt.ylabel('True label')
plt.xlabel('Predicted label')
```

```
[213]: Text(0.5, 47.722222222222, 'Predicted label')
```



```
[145]: # #testing 2
# class Shallow_ANN2():

# def fit(self, X, y, neurons=6, eta=1e-3, epochs=1e3, show_curve= True):
```

```

#     epochs= int(epochs)
#     N, D = X.shape
#     Y = one_hot_encode(y)
#     K= Y.shape[1]

#     #Weights Initialization
#     self.W =[l: np.random.randn(M[0], M[1]) for l, M in enumerate(zip([D,
↪neurons],[neurons, K]), 1)}
#     self.B = {l: np.random.randn(M) for l, M in enumerate([neurons, K], 1)}
#     #Define Activations
#     self.a ={1:np.tanh, 2:softmax}

#     J= np.zeros(epochs)

#     #SGD Steps
#     for epoch in range(epochs):
#         self.__forward__(X)
#         J[epoch]= cross_entropy(Y, self.Z[2])
#         #Weight update rules for output layer(Layer 2)
#         self.W[2] -= eta*(1/N)* self.Z[1].T@(self.Z[2] - Y)
#         self.B[2] -= eta*(1/N)*(self.Z[2]-Y).sum(axis=0)
#         #Weight Update Rule for Layer 1
#         #self.W[1] -= eta*(1/N)*X.T@((self.Z[2]-Y)@self.W[2].T*(1-self.Z[1]**2))
#         #self.B[1] -= eta*(1/N)*((self.Z[2]-Y)@self.W[2].T*(1-self.Z[1]**2)).
↪sum(axis=0)
#         self.W[1] -= eta*(1/N)*X.T@((self.Z[2]-Y)@self.W[2].T*(1-self.Z[1]**2))
#         self.B[1] -= eta*(1/N)*((self.Z[2]-Y)@self.W[2].T*(1-self.Z[1]**2)).
↪sum(axis=0)

#     if show_curve:
#         plt.figure()
#         plt.plot(J)
#         plt.xlabel('epochs')
#         plt.ylabel('$\mathcal{J}$')
#         plt.title('Training Curve')
#         plt.show()

#     def __forward__(self, X):
#         self.Z ={0:X}
#         for l in sorted(self.W.keys()):
#             self.Z[l] = self.a[l](self.Z[l-1]@self.W[l] + self.B[l])

#     def predict(self, X):
#         self.__forward__(X)
#         return self.Z[2].argmax(axis=1)

```

```
[146]: # my_ann2 = Shallow_ANN2()
# my_ann2.fit(X_train, y_train, neurons= 2, eta=0.001, epochs = 30000,
↳ show_curve= True )
# y_hat_my_ann2 =my_ann2.predict(X_test)

# print("Accuracy: ", accuracy(y_test, y_hat_my_ann2))
# print(my_ann2.W)
# print(my_ann2.B)
```

```
[147]: # #neurons_list = [2, 3, 5]
# # eta_list = [1e-3, 1e-2, 1e-1]
# # epochs_list = [1e4, 2e4, 3e4]

# neurons_list = [2, 3, 4]      # Slight increase in neuron count
# eta_list = [0.001, 0.002, 0.005] # Fine-tuning the learning rate
# epochs_list = [20000, 25000, 30000] # Reducing epochs
# l2_list = [0.001, 0.01]

# best_accuracy = 0
# best_params = {'neurons': None, 'eta': None, 'epochs': None}

# # Loop through different combinations of hyperparameters
# for neurons in neurons_list:
#     for eta in eta_list:
#         for epochs in epochs_list:
#             my_ann = Shallow_ANN() # Initialize your ANN
#             my_ann.fit(X_train, y_train, neurons=neurons, eta=eta,
↳ epochs=int(epochs), show_curve=False)

#             # Predict and calculate accuracy
#             y_hat_2LFF = my_ann.predict(X_test)
#             accuracy_val = accuracy(y_test, y_hat_2LFF)

#             # Keep track of the best accuracy and parameters
#             if accuracy_val > best_accuracy:
#                 best_accuracy = accuracy_val
#                 best_params['neurons'] = neurons
#                 best_params['eta'] = eta
#                 best_params['epochs'] = epochs

# print(f"Best Accuracy: {best_accuracy:.4f}")
# print(f"Best Parameters: Neurons={best_params['neurons']},
↳ Eta={best_params['eta']}, Epochs={best_params['epochs']}")
```

4 Artificial Neural Net with Back Propagation and Variable Architecture

```
[214]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
[215]: # Useful functions
# Activations
def linear(H):
    return H

def ReLU(H):
    return H*(H>0)

def sigmoid(H):
    return 1/(1+np.exp(-H))

def softmax(H):
    eH = np.exp(H)
    return eH/eH.sum(axis=1, keepdims= True)

# Loss Functions

def cross_entropy(Y, P_hat):
    return -(1/len(Y))*np.sum(Y*np.log(P_hat)))

def binary_cross_entropy(y, p_hat):
    return -(1/len(y))*np.sum(y*np.log(p_hat)+ (1-y)*np.log(1-p_hat))

def OLS(Y, Y_hat):
    return (1/(2*len(Y)))* np.sum((Y-Y_hat)**2)

#Metrics
def accuracy(y, y_hat):
    return np.mean(y == y_hat)

def R2(y, y_hat):
    return 1-np.sum((y-y_hat)**2)/ np.sum((y-y.mean())**2)

# Misc
def one_hot(y):
    N= len(y)
    K=len(set(y))
    Y= np.zeros((N,K))
    for i in range(N):
```

```

    Y[i,y[i]]=1
    return Y

```

[216]: *# Derivatives of Activation Functions*

```

def derivative(Z,a):

    if a==linear:
        return 1

    elif a==sigmoid:
        return Z*(1-Z)

    elif a==np.tanh:
        return 1-Z**2

    elif a==ReLU:
        return (Z>0).astype(int)

    else:
        ValueError("Unknown Activation")

```

[151]: *# # Class*

```

# class ANN():

#     def __init__(self, architecture, activations=None, mode=0):
#         self.mode=mode
#         self.architecture=architecture
#         self.activations= activations
#         self.L = len(architecture)+1

#     def fit(self, X, y, eta=1e-3, epochs=1e3, show_curve=False):
#         epochs=int(epochs)
#         #Classifier, mode=0, Regressor, mode=1
#         if self.mode:
#             Y=y
#             K=1
#         else:
#             #Y =one_hot(y)
#             Y=y.reshape(-1,1)
#             #Y=y
#             K =Y.shape[1]

#         N,D = X.shape

#         #Iniatize Weights(and Biases)
#         self.W = {l: np.random.randn(M[0],M[1]) for l, M in
# ↪ enumerate(zip((D)+self.architecture), (self.architecture+[K]),1))}

```

```

#     self.B = {l: np.random.randn(M) for l,M in enumerate(self.
↪architecture+[K],1)}

#     #Activation Setup
#     if self.activations is None:
#         self.a = {l:ReLU for l in range(1,self.L)}
#     else:
#         self.a = {l: act for l,act in enumerate(self.activations,1)}

#     #Output activation Functions
#     if self.mode:
#         self.a[self.L]=linear
#     else:
#         self.a[self.L]=sigmoid

#     #Define Loss
#     J = np.zeros(epochs)

#     #Training Cycle
#     for epoch in range(epochs):
#         self.__forward__(X)

#         if self.mode:
#             J[epoch]= OLS(Y, self.Z[self.L])
#         else:
#             J[epoch]=binary_cross_entropy(Y, self.Z[self.L])
#         #Back Prop
#         dH = (1/N)*(self.Z[self.L]-Y)
#         for l in sorted(self.W.keys(), reverse=True):

#             dW = self.Z[l-1].T@dH
#             dB = dH.sum(axis=0)
#             #Weight Update Rules per layer
#             self.W[l] -=eta*dW
#             self.B[l] -=eta*dB

#         if l>1:
#             dZ =dH@self.W[l].T
#             dH = dZ*derivative(self.Z[l-1],self.a[l-1])

#     if show_curve:
#         plt.figure()
#         plt.plot(J)
#         plt.xlabel("epochs")
#         plt.ylabel("$\mathcal{J}$")
#         plt.title("Training Curve")
#         plt.show()

```

```

# def __forward__(self, X):
#     self.Z = {0:X}
#     for l in sorted(self.W.keys()):
#         self.Z[l] = self.a[l](self.Z[l-1]@self.W[l]+self.B[l])

# def predict(self, X):
#     self.__forward__(X)
#     if self.mode:
#         return self.Z[self.L]
#     else:
#         #return self.Z[self.L].argmax(axis=1)
#         return (self.Z[self.L] > 0.5).astype(int)

```

```

[152]: # import numpy as np

# # Define ranges for hyperparameters
# architectures = [[6, 10, 4], [6, 12, 4]] # Different architectures to try
# eta_list = [1e-4, 1e-3, 1e-2] # Learning rate options
# epochs_list = [50000, 100000, 150000] # Epoch options
# l2_list = [0.001, 0.01] # Regularization strength options
# dropout_list = [0.2, 0.3] # Dropout rates

# best_accuracy = 0
# best_params = {'architecture': None, 'eta': None, 'epochs': None, 'l2': None,
#               ↪ 'dropout': None}

# # Loop through all combinations of hyperparameters
# for architecture in architectures:
#     for eta in eta_list:
#         for epochs in epochs_list:
#             for l2 in l2_list:
#                 for dropout_rate in dropout_list:
#                     # Initialize the ANN model with the current architecture,
#                     ↪ L2 regularization, and dropout
#                     my_ann_classifier = ANN(architecture=architecture,
#                     ↪ activations=[np.tanh] *
#                     ↪ len(architecture),
#                     ↪ l2=l2, dropout_rate=dropout_rate)

#                     # Fit the model with the current hyperparameters
#                     for epoch in range(int(epochs)):
#                         my_ann_classifier.fit(X_train, y_train, eta=eta,
#                         ↪ epochs=1, show_curve=False) # Training for 1 epoch at a time

#                         # Predict on the test set and calculate accuracy
#                         ↪ after each epoch

```



```

#                                     y_hat_ANN = my_ann_classifier.predict(X_test)
#                                     accuracy_val = accuracy(y_test, y_hat_ANN)

#                                     # Print accuracy and parameters for each epoch
#                                     print(f"Epoch: {epoch + 1}, Architecture:␣
↳{architecture}, Eta: {eta}, L2: {l2}, Dropout: {dropout_rate}, Accuracy:␣
↳{accuracy_val:.4f}")

#                                     # Track the best accuracy and the best parameters
#                                     if accuracy_val > best_accuracy:
#                                     best_accuracy = accuracy_val
#                                     best_params['architecture'] = architecture
#                                     best_params['eta'] = eta
#                                     best_params['epochs'] = epoch + 1 # Track the␣
↳epoch at which the best result occurred
#                                     best_params['l2'] = l2
#                                     best_params['dropout'] = dropout_rate

# # Print the best results
# print(f"Best Accuracy: {best_accuracy:.4f}")
# print(f"Best Parameters: Architecture={best_params['architecture']},␣
↳Eta={best_params['eta']}, "
#     f"Epochs={best_params['epochs']}, L2={best_params['l2']},␣
↳Dropout={best_params['dropout']}")

```

```

[153]: # my_ann_classifier = ANN(architecture=[6,8, 4], activations=[np.tanh]*3)
# my_ann_classifier.fit(X_train,y_train, eta=1e-3, epochs=1e5, show_curve=True)
# y_hat_ANN=my_ann_classifier.predict(X_test)

# print(my_ann_classifier.W)
# print(my_ann_classifier.B)
# print(f"Testing Accuracy: {accuracy(y_test,y_hat_ANN):0.4f}")

```

[153]:

```

[217]: class ANN():

    def __init__(self, architecture, activations=None, mode=0):
        self.mode=mode
        self.architecture=architecture
        self.activations= activations
        self.L = len(architecture)+1

    def fit(self, X, y, eta=1e-3, epochs=1e3, show_curve=False):
        epochs=int(epochs)
        #Classifier, mode=0, Regressor, mode=1
        if self.mode:

```

```

        Y=y
        K=1
    else:
        Y =one_hot(y)
        #Y = y.reshape(-1, 1)
        K =Y.shape[1]

N,D = X.shape

#Initiatize Weights(and Biases)
self.W = {l: np.random.randn(M[0],M[1]) for l, M in enumerate(zip(([D]+self.
↪architecture), (self.architecture+[K])),1))}
self.B = {l: np.random.randn(M) for l,M in enumerate(self.
↪architecture+[K],1)}

#Activation Setup
if self.activations is None:
    self.a ={l:ReLU for l in range(1,self.L)}
else:
    self.a = {l: act for l,act in enumerate(self.activations,1)}

#Output activation Functions
if self.mode:
    self.a[self.L]=linear
else:
    self.a[self.L]=softmax

#Define Loss
J = np.zeros(epochs)

#Training Cycle
for epoch in range(epochs):
    self.__forward__(X)

    if self.mode:
        J[epoch]= OLS(Y, self.Z[self.L])
    else:
        J[epoch]=cross_entropy(Y, self.Z[self.L])

    #Back Prop
    dH = (1/N)*(self.Z[self.L]-Y)
    for l in sorted(self.W.keys(), reverse=True):

        dW = self.Z[l-1].T@dH
        dB = dH.sum(axis=0)
        #Weight Update Rules per layer
        self.W[l] -=eta*dW
        self.B[l] -=eta*dB

```

```

        if l>1:
            dZ =dH@self.W[l].T
            dH = dZ*derivative(self.Z[l-1],self.a[l-1])

    if show_curve:
        plt.figure()
        plt.plot(J)
        plt.xlabel("epochs")
        plt.ylabel(" $J$ ")
        plt.title("Training Curve")
        plt.show()

    def __forward__(self, X):
        self.Z ={0:X}
        for l in sorted(self.W.keys()):
            self.Z[l] = self.a[l](self.Z[l-1]@self.W[l]+self.B[l])

    def predict(self, X):
        self.__forward__(X)
        if self.mode:
            return self.Z[self.L]
        else:
            return self.Z[self.L].argmax(axis=1)
            #return (self.Z[2] > 0.5).astype(int)

```

```

[218]: architecture_list = [[6, 8, 4], [6, 10, 4], [6, 12, 6]]
eta_list = [0.01, 0.05, 0.1]
epochs_list = [2000, 4000, 6000]

best_accuracy = 0
best_params = {'architecture': None, 'eta': None, 'epochs': None}

for architecture in architecture_list:
    for eta in eta_list:
        for epochs in epochs_list:
            print(f"Trying: Architecture={architecture}, Eta={eta},
↳Epochs={epochs}")
            my_ann_classifier = ANN(architecture=architecture, activations=[np.
↳tanh] * len(architecture))
            my_ann_classifier.fit(X_train, y_train, eta=eta, epochs=epochs,
↳show_curve=False)

            y_hat_ANN = my_ann_classifier.predict(X_test)
            accuracy_val = accuracy(y_test, y_hat_ANN)
            print(f"Accuracy: {accuracy_val:.4f} | Architecture:
↳{architecture}, Eta: {eta}, Epochs: {epochs}")

```

```

        if accuracy_val > best_accuracy:
            best_accuracy = accuracy_val
            best_params['architecture'] = architecture
            best_params['eta'] = eta
            best_params['epochs'] = epochs

print(f"Best Accuracy: {best_accuracy:.4f}")
print(f"Best Parameters: Architecture={best_params['architecture']},  

      ↪Eta={best_params['eta']}, Epochs={best_params['epochs']}")


```

```

[221]: my_ann_classifier = ANN(architecture=[6,8, 4], activations=[np.tanh]*3)
#neurons= 2,  eta=0.001, epochs = 30000
#epochs= 4000, eta= 0.1
#neurons= 4,  eta=0.1, epochs = 8000
#[6, 8, 4], Eta=0.05, Epochs=2000 ### Good Good
my_ann_classifier.fit(X_train,y_train, eta=0.1, epochs=8000, show_curve=True)
y_hat_ANN=my_ann_classifier.predict(X_test)

print(my_ann_classifier.W)
print(my_ann_classifier.B)
print(f"Testing Accuracy: {accuracy(y_test,y_hat_ANN):0.4f}")

```



```

{1: array([[ 1.39885940e-01, -3.46816181e-01, -2.44269957e-03,
           -1.46614427e+00,  4.54441082e-01, -8.26840465e-02],
          [-2.22183142e-01,  4.88728903e-01,  7.27213154e-01,
           -1.12152277e+00, -1.31947946e+00,  1.02520541e+00],
          [ 1.09709879e+00,  2.44498451e+00,  5.97530229e-02,
           -1.20171851e+00,  9.60928561e-02, -1.12679577e-01],
          [-6.53403824e-01, -3.79915557e-01,  4.98311993e-01,
           -9.65117650e-01,  8.83925045e-01, -4.81400196e-01],
          [-5.77344137e-01, -4.78135093e-01,  1.76498100e+00,
           -8.93951301e-01,  5.58120743e-01, -2.90120389e+00],
          [ 5.70413670e-01,  1.59835684e+00,  9.78115527e-03,
           6.62456461e-01, -1.24701471e+00, -1.46914999e-02],
          [-8.97547730e-01,  1.16012852e+00, -2.75756094e-01,
           -6.01603914e-01,  7.14411339e-01, -1.51353098e-01],
          [-3.44007867e-01,  1.16650798e+00, -6.35001654e-02,
           4.71958758e-01, -4.80404889e-01,  6.21837697e-02],
          [-1.34290954e+00,  7.40478624e-01, -2.01108437e-01,
           1.23856882e-02, -1.12685441e+00,  5.03098793e-01],
          [ 7.19026458e-01, -3.24808592e-01,  1.92167047e-01,
           -9.72255360e-01, -1.81734376e+00, -1.14926915e-01],
          [ 1.09411049e+00,  1.87023962e+00,  1.11856275e-01,
           -1.57432674e+00, -1.57010973e+00, -2.23657156e-01]]), 2: array([[
2.54173771e-01,  4.18619101e-01, -3.66365072e-01,
           -4.95916488e-01, -1.37515834e+00, -2.76525393e-03,
           5.96645601e-01, -2.34943298e-01],
          [-1.17973680e+00,  1.48744195e-01,  4.01671401e-01,
           -2.24036882e+00,  7.81367073e-01, -7.78252378e-01,
           -8.46743477e-01,  1.19404024e+00],
          [-1.67444299e+00, -6.68664425e-01,  1.77160215e+00,
           9.18652919e-01,  3.19068975e-01,  4.82779473e-01,
           -4.32323611e-01, -1.90636568e+00],
          [-3.33239476e-01,  3.21441233e-01, -7.11006115e-01,
           -3.56910607e-02, -1.85439018e+00,  4.73392543e-01,
           -5.69565800e-01,  1.66077869e+00],
          [-7.68647032e-01, -6.95556368e-01,  8.47658157e-01,
           8.57749603e-01, -4.31326891e-01, -2.59540975e+00,
           1.92448745e+00, -1.55932690e-01],
          [ 6.98537002e-01,  3.49419916e+00,  1.70145171e-01,
           -4.66692763e-01, -5.27498013e-01, -6.26288308e-01,
           -1.13384006e+00, -2.43869150e+00]]), 3: array([[ -1.69736097,
1.78100354,  0.17887751,  0.53644517],
          [ 0.3459838 ,  0.36028619,  1.16412572,  1.93996936],
          [ 1.17986617,  0.55335737,  0.4335691 ,  0.022611 ],
          [-1.84208271, -0.37581754, -0.12525235,  0.90426299],
          [-0.44722313,  0.157237 , -0.01899177, -0.32723522],
          [ 0.33182308, -0.29660853, -0.3262923 ,  1.19500884],

```

```

        [-0.82534787,  0.59978189,  0.90385401,  1.88020518],
        [-0.96620902,  0.26941558, -1.35628811,  2.75031284]]), 4:
array([[ -2.70499896,  0.74720213],
       [ 0.50182    , -0.77742429],
       [-0.98312157,  0.33932148],
       [ 1.28053735,  0.23166609]])}
{1: array([-0.66249792,  0.79834909, -1.73961982, -0.69830509, -0.63001626,
          0.21072266]), 2: array([ 0.85803675,  0.62628061, -0.31820561,
          0.30859863, -0.67971461,
          1.34473707, -0.05289225,  0.32356479]), 3: array([ 0.95387333,
          0.40643579, -0.76911796,  0.11179271]), 4: array([-1.61178633,  0.38996769])}
Testing Accuracy:  0.7796

```

```


[222]: ## Artificial Neural Net with Back Propagation
plt.figure(figsize=(10,7))
y_actual_ANN = pd.Series(y_test, name='Actual')
y_pred_ANN = pd.Series(y_hat_ANN, name='Predicted')
cm = pd.crosstab(y_actual_ANN, y_pred_ANN)
ax = sns.heatmap(cm, annot=True, fmt="d")
plt.title("Artificial Neural Net with Back Propagation Confusion Matrix")
plt.ylabel('True label')
plt.xlabel('Predicted label')

```

```

[222]: Text(0.5, 47.72222222222222, 'Predicted label')

```



```
[ ]: # my_ann_classifier = ANN(architecture=[6,8, 4], activations=[ReLU] * (2) + [np.
      ↪ tanh])
      # my_ann_classifier.fit(X_train,y_train, eta=0.05, epochs = 6000, ↪
      ↪ show_curve=True)
      # y_hat_ANN=my_ann_classifier.predict(X_test)

      # print(my_ann_classifier.W)
      # print(my_ann_classifier.B)
      # print(f"Testing Accuracy: {accuracy(y_test,y_hat_ANN):0.4f}")
```

```
[ ]: # def scale_new_data_point(new_data_point, scaling_params, columns):
      #     scaled_data_point = []
      #     for i, column in enumerate(columns):
      #         if column in scaling_params:
      #             min_value = scaling_params[column]['min']
      #             max_value = scaling_params[column]['max']
      #             scaled_value = (new_data_point[i] - min_value) / (max_value - ↪
      ↪ min_value)
      #             scaled_data_point.append(scaled_value)
```

```

#         else:
#             scaled_data_point.append(new_data_point[i])

#     return scaled_data_point

# columns = ['longitude', 'latitude', 'lot_acres', 'sqrt_ft', 'bedrooms',
# ↪ 'bathrooms', 'garage', 'fireplaces']
# teacher_test = [[-110.3782, 31.356362, 2154, 10500, 13, 10, 0, 6]]

# scaled_teacher_test = [scale_new_data_point(point, scaling_params, columns)
# ↪ for point in teacher_test]
# scaled_teacher_test

# Test_case = ANN.predict(scaled_teacher_test)
# Test_case

```

[]: